

Fairer access starts with the right target.

Historical awards measure the committee. They do not establish who should receive support.

4,000

Applicants scored

All evaluation IDs preserved

1,600

Grants allocated

Exact 40% budget

0.021 pp

Regional access gap

Selection-rate gap; not equal opportunity

Our contribution: diagnose the proxy problem, repair the ranking criteria, and enforce an auditable allocation budget.

Independent-reference accuracy is unknown. A 98% accuracy claim is not supported.

Deleting the region column leaves the gap.

Same 3,000-row holdout, seed 42. Metrics below concern historical committee decisions.

18.76 pp

Full baseline

88.13% historical agreement

18.05 pp

Without region

Postal and other proxies remain

17.31 pp

Without region + postal

Distance, income and hours carry signal

Historical grant rates: 48.4% in centres versus 27.3% in remote regions. Regional differences persist within academic-score bands.

Distance alone predicts regional block with ROC AUC 0.997 on held-out data. This identifies a proxy; it does not establish a causal effect.

Repair the score. Make the allocation explicit.

Policy choices are disclosed; no synthetic labels are passed off as independent merit.

1

Estimate conditional associations

Fit committee decisions while controlling for geography, wealth, program and other observed attributes.

2

Apply the approved direction of effects

Remove location/program effects, positive wealth advantage and first-generation penalty. Keep nonnegative academic and work effects.


3

Solve the exact allocation problem

Take the best within-group scores under 40% total grants and a maximum 2-point selection-rate gap. Stable tie breaking; no outcome labels required.

Show the trade-off, including its limits.

12 constraint settings per score, each with an exact 40% budget. Star = chosen policy.



Historical agreement is a diagnostic comparator. Hidden-reference utility and opportunity remain unknown.

A valid, reproducible allocation.

The score repair already produces near parity; the 2-point guardrail is inactive on the final cohort.

40.01%

Centre selection

949 grants / 2,372 applicants

39.99%

Remote selection

651 grants / 1,628 applicants

40.00%

Overall budget

Inside the official 36%-44% range

Verified: exact IDs and row count, binary decisions, deterministic rerun, monotonic score rules, and allocator optimality against exhaustive small cases.

Held-out historical agreement: 85.33%. Historical-label TPR gap increases from 0.063 to 0.169; this does not determine the hidden-merit result.

Govern decisions beyond the leaderboard.

Named responsibilities, measurable triggers, a human appeal and a controlled fallback.

Independent reference

Regional panel reviews funded and rejected applicants under a published eligibility rubric. Validate true-merit TPR and utility before release.

Every batch

ML owner blocks schema or budget failures. Monitor regional access, intersections, input drift and model versions with denominators.

Appeal and correction

Human officer reviews errors and context. Explain the score and regional constraint. Reconcile funding through an approved budget process.

Freeze and reassess

Escalate independently measured opportunity gaps. Freeze automatic releases on severe incidents and use the approved human process.

Approve the criteria. Measure independent merit.

The hidden reference cannot be recovered merely by fitting historical decisions more accurately.

Known

A complete audit, exact-budget predictions, an inspectable model, a constraint sweep and a monitoring plan.

Still to establish

True-merit accuracy and equal opportunity. Appropriate work and financial-need weights. Applicant-level consequences.

Our standard: transparent assumptions, measured evidence, and a path to challenge a decision.

Sources: supplied IVADO briefs, HxBuddy track, Fairlearn docs. Tools and Codex assistance disclosed in README.